

tfConstrainedGauss Python package

Oliver K. Ernst

September 25, 2021

This package implements two methods for finding a sparse precision matrix with a given structure from a given covariance matrix.

1 Identity-based method

Given an $n \times n$ covariance matrix, here of size $n = 3$:

$$\Sigma = \begin{pmatrix} c_{11} & c_{12} & c_{13} \\ c_{12} & c_{22} & c_{23} \\ c_{13} & c_{23} & c_{33} \end{pmatrix} \quad (1)$$

and given the structure of the precision matrix (i.e. given the Gaussian graphical model), for example:

$$P = \begin{pmatrix} p_{11} & p_{12} & 0 \\ p_{12} & p_{22} & p_{23} \\ 0 & p_{23} & p_{33} \end{pmatrix} \quad (2)$$

(note that the diagonal elements are always non-zero), the goal is to find the elements of the precision matrix by:

$$P^* = \underset{P}{\operatorname{argmin}} |P\Sigma - I| \quad (3)$$

where I is the identity.

The advantage of this approach is that it does not require calculating the inverse of any matrix, particularly important for large n .

The disadvantage of this approach is that the solution found for P may not yield a covariance matrix P^{-1} whose individual elements are close to those of Σ . That is, while $P\Sigma$ may be close to the identity, there are likely errors in every single element of P^{-1} .

2 MaxEnt-based method

Given the structure of the $n \times n$ precision matrix (i.e. given the Gaussian graphical model), for example:

$$P = \begin{pmatrix} p_{11} & p_{12} & 0 \\ p_{12} & p_{22} & p_{23} \\ 0 & p_{23} & p_{33} \end{pmatrix} \quad (4)$$

(note that the diagonal elements are always non-zero), and given the covariances for corresponding to every *non-zero* entry in P , i.e. given:

$$c_{11}, c_{12}, c_{22}, c_{23}, c_{33} \quad (5)$$

the goal is to find the elements of P . In other words, every unique element (i, j) of the $n \times n$ symmetric matrix has a given constraint, either to a value in the covariance matrix, or a zero entry in the precision matrix.

This is a maximum entropy (MaxEnt) setup. The elements of the precision matrix p_{ij} are directly the interactions in the Gaussian graphical model; the moments they control in a MaxEnt sense are the covariances c_{ij} .

The problem can be solved in a number of ways, for example using Boltzmann machine learning, where we minimize:

$$P^* = \operatorname{argmin}_P \mathcal{D}_{\mathcal{KL}} = \min_P \sum_n p(n) \ln \frac{p(n)}{\tilde{p}(n)} \quad (6)$$

where $p(n)$ is the (unknown) data distribution that gave rise to the given covariances c_{ij} and $\tilde{p}(n)$ is the Gaussian with precision matrix P . The gradients that result are the wake sleep phase:

$$\Delta p_{ij} \propto c_{ij} - (P^{-1})_{ij} \quad (7)$$

In TensorFlow, we minimize the MSE loss for the individual terms, which results in the same first order gradients:

$$P^* = \operatorname{argmin}_P \sum_{ij} \left\| c_{ij} - (P^{-1})_{ij} \right\|_2 \quad (8)$$

To learn each element of the covariance matrix with equal importance, we can use a weighted MSE loss:

$$P^* = \operatorname{argmin}_P \sum_{ij} w_{ij} \left\| c_{ij} - (P^{-1})_{ij} \right\|_2 \quad (9)$$

where

$$w_{ij} = c_{ij}^{-2} \quad (10)$$